

Technotes on LatSeq tool

gitlab.eurecom.fr/oai/openairinterface5g

F. Ronteix-Jacquet

flavien.ronteixjacquet@orange.com



Introduction

Measurement module

Analysis Module

Example

Scenarios

Conclusion

Introduction I

LatSeq for *Latency Sequence Analysis* is a tool developped by A. Ferrieux and F. Ronteix-Jacquet dM at Orange Labs Networks to analyse internal latency. This tool is designed to be agnostic to the system analyzed but, first developed for OAI's base station.

It is composed of 2 parts, a *measurement* module and an *analysis* module. The architecture in figure below :

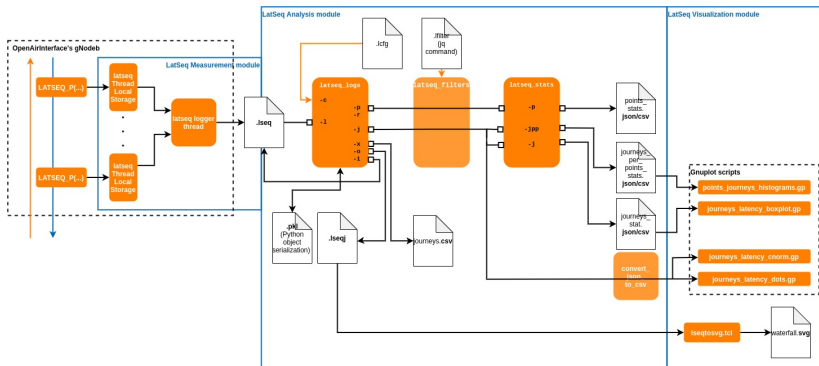


Figure – LatSeq tool architecture

Introduction II

The normal usage of LatSeq is as follow :

- ▶ run a **OAI RAN** on a scenario
- ▶ log data at **measurement points** in the BS (`LATSEQ_P(protocol.ent.func, format_for_dataid, dataid)`)
- ▶ export **log.lseq** file
- ▶ parsing by **latseq_logs** script
- ▶ rebuild packets' journey, possible paths,...
- ▶ filter logs with **latseq_filter**
- ▶ compute stats and display **visualization** (latency sequences, latency stats,...)

Measurement module I

LatSeq Architecture Figure1 shows that a part of LatSeq code is implemented in OAI source code. This code inside OAI code is composed of

- ▶ *latseq logger thread* which is relatively independant from OAI (but must to be started by the main thread)
- ▶ *latseq points* (LATSEQ_P(...)) inside OAI's functions. This part is really dependent of OAI Code... The knowledge of how OAI works is really important to have meaningfull measurements.

In source code :

- ▶ source directory for LatSeq measurement module : `common/UTILS/LATSEQ`
- ▶ source directory for LatSeq analysis module : `common/UTILS/LATSEQ/tools`
- ▶ build_oai flag : `-enable-latseq`
- ▶ macro enable : `LATSEQ`
- ▶ macro to log : `LATSEQ_P(point_identifier, data_identifier_format, data_identifiers...)`. ex. `LATSEQ_P("Dpdcp.tx-rlc.tx.um", "len%d:rnti%d:drb%d:psn%d:lcid%d:rsdu%d", ssize, ctxt_pP->rnti, rlc_p->rb_id, seqnum, rlc_p->channel_id, ssize);`

Measurement module II

How it works in OAI?

- ▶ global logger_thread structure *g_latseq* (see Figure3)
 - ▶ flags : *is_running*, *is_debug*
 - ▶ init internals : *filelog_name*, *time_zero*, *rdtsc_zero*, *FILE * outstream*, *latseq_stats*
 - ▶ *log_buffer* : buffer of *latseq_element* to record, behave like a circular array
 - ▶ writer's and reader's positions in the buffer
- ▶ at the init, *init_latseq()* to call : init global latseq context and call *init_logger_latseq()* which init logger thread (see Figure2)
- ▶ the logger thread is a C's pthread which has the task to write elements from the buffer into log file. This needs to be in a different process, in background to avoid undesirable latency in logging. It is running until *is_running* flag is not 0 or until there is an error in writing files or until the log buffer is full.
- ▶ *LATSEQ_P(...)* is a macro which calls *_write_latseq_entry()*. This function try to add fastly a new measurement in the *log_buffer*. Time constraint on this function is huge because it runs in the oai threads.
- ▶ The buffer is a circular buffer which means, the writer's head should always be in advance of reader's head. If the buffer is almost full, we prefer to stop immediately latseq and print an error.

Measurement module III

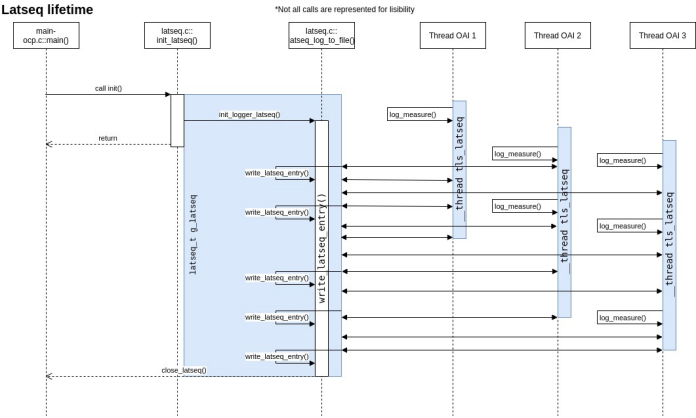


Figure – LatSeq Lifetime

Measurement module IV

LatSeq data structures

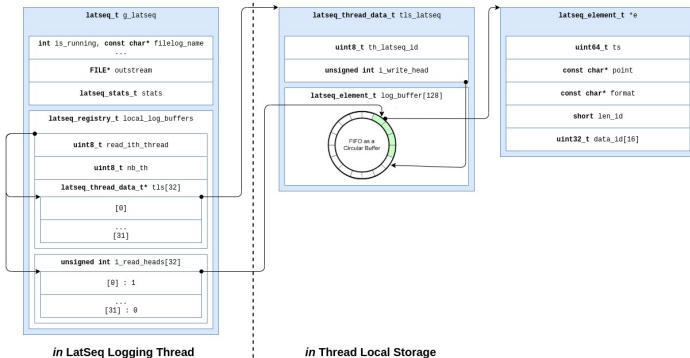


Figure – Latseq data structures

To ensure the objective of low-dependency, low-impact, low-memory requirements for measurements, these features has been implemented :

- ▶ Usage of `rdtsc` asm instruction to take timestamp
- ▶ Expansion, conversion and save of the entries **asynchronously** by the `logger_thread` and not by the `oai_thread` which performs the measurement
- ▶ Use of **inline function** for `log_measure()`

Measurement module V

- ▶ Use of **thread-local storage** (TLS) for thread-local buffer for measures instead of one big buffer in logger thread with the problem of mutexes, locks and waits. Thread local storage allows to reserve memory for LatSeq for all threads with a measurement point.
- ▶ Use of **lock-less ring buffer** for measurements.
- ▶ **Parametric Macro** instead of variadic function for `log_measure`. For 1 to 10 data identifiers, only one macro to use, `LATSEQ_P(...)`. We try to use a variadic function to have only one function, however, after test, this solution introduce a big overhead on instruction (11 to 19ns for 1 identifier, 16.3ns to 61ns for 10 identifiers).
- ▶ use of **const char*** for latseq point's name in `LATSEQ_P`. The copy takes only few asm instructions to be performed. Most efficient than allocates memory and free it later.
- ▶ Use of **fprintf()** for swift write of logging in file (with its buffering property)

Additional notes :

Measurement module VI

- ▶ The ring buffers size are set to **32** latseq logs elements. The buffer is filled by the LATSEQ_P points of the thread. If the thread has several Latseq measurement points, then this point log in the same ring buffer. Ring buffers are emptied by the logger_thread. Logger_thread visit all ring buffer of registered threads in turn. If the selectionned buffer is empty, then go to the next directly. By this, the most active ring buffer are emptied regularly. The ring buffer has to be sized according to the most active thread (in term of measurements). By experimentations, the speed of emptying is mostly conditionned by the disk speed. Then, the speed of measurement should not be to large of disk speed for a too long time. Buffers absorb this differences and absorb the time to visit all thread's buffers.
- ▶ To ensure we can rebuild packet's path and perform statistics and so on, at least one data identifier should link 2 measurements points. More local data identifier is given, less there are ambiguity.
- ▶ to ensure the journeys are rebuildable, the measurement points installer ought to be careful with : 1) the existence of a path between an In and an Out point (path existence property), 2) the global identifier never changes through the path of one journey (unique global identifier property), 3) at least one local identifier links 2 connected measurement points (the continuity property), 4) avoir loop between uplink and downlink's paths (strict direction separation property)

the measurement points installer ought to be careful with :

- ▶ The existence of a path between an *Input* and *Output* point (*Path existence property*).
- ▶ The global identifier is consistent throughout the journey's path (*Unique global identifier property*).
- ▶ At least one local identifier links 2 connected points (*The continuity property*).
- ▶ No bridges between Uplink and Downlink's paths (*Strict directions separation property*).

Analysis Module I

How the post-processing and analysis module works ?

- ▶ from lseq logs, extract and clean measurements
- ▶ rebuild paths with an algorithm
- ▶ **latseq_logs** script could be called with this options (takes *.lseq in input and return *.json) :
 - ▶ -i : returns the inputs after filtering and cleaning as a string
 - ▶ -r : returns the paths present in the log file as json.
 - ▶ -p : returns points structure as json. (Becareful, if journeys has not been rebuilt, then you do not have "duration" attribute which is used for statistics)
 - ▶ -j : returns journeys structure as json. This is the main options to have statistics on packet's latency
 - ▶ -o : returns a latseq journey file line by line. redirects output to a file to have a *.lseqj and build "waterfall" visualization
- ▶ filter with **latseq_filter** to get statistics on specific category. for instances only downlinks points or packets with drb==1.
- ▶ request statistics to **latseq_stats** with options : (takes json in input and returns json or text) :
 - ▶ -j : returns statistics on journeys
 - ▶ -p : returns statistics on points

Example I

An example of LatSeq implementation for OAI LTE's eNB. The scenario is soft-NodeB with a LTE core. The UE send ping (ICMP) to a target.

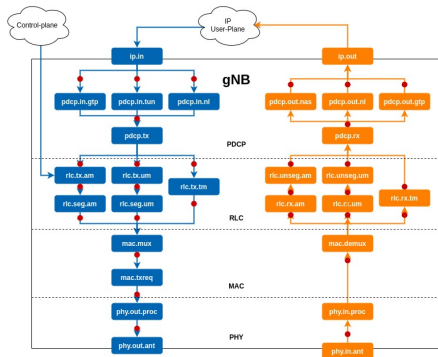


Figure – Latseq data points in OAI's eNB source code

Example II

```
1587645146.191801 D ip--pdcp.in.gtp len64:rnti54614:drb1.gsn12
1587645146.191802 D pdcp.in--pdcp.tx len64:rnti54614:drb1.gsn12.psn10
1587645146.191803 D pdcp.tx--rlc.tx.um len66:rnti54614:drb1.psn10.lcid3.rsdu0
1587645146.191962 D rlc.tx.um--rlc.seg.um len66:rnti54614:drb1.lcid3.rsdu0.rsn7.rso0
1587645146.191962 D rlc.seg.um--mac.mux len68:rnti54614:drb1.lcid3.rsn7.fm899
1587645146.191965 D mac.mux--mac.txreq len73:rnti54614:lcid3.harq7.txreq0.fm899.subfm1
1587645146.191997 D mac.txreq--phy.out.proc len73:rnti54614:ue2.harq7.fm899.subfm1
1587645146.207988 U phy.in.proc--mac.demux len60:rnti54614:lcid3.ue0.fm900
1587645146.207989 U mac.demux--rlc.rx.um len60:rnti54614:drb1.lcid3.rsn16.fm900
1587645146.207989 U rlc.rx.um--rlc.unseg.um len58:rnti54614:drb1.lcid3.rsn16.fm900
1587645146.207990 U rlc.unseg.um--pdcp.rx len54:rnti54614:drb1.lcid3.rsn16.psn0.psn12.fm900
1587645146.207994 U pdcp.rx--ip len52:rnti54614:drb1.psn12.fm900
```

Figure – Example of Latseq Logs at the output of Measurement module

Example of scripts pipeline

```
> ./latseq_logs.py -p -l ocp-softmodem.lseq | \
./latseq_filter.sh points_downlinks.lfilter | \
./latseq_stats.py -sp -P
> ./latseq_logs.py -j -l ocp-softmodem.lseq | \
./latseq_stats.py -f csv -djd > durations.csv | \
./gnuplot -e "data='durations.csv'" journeys_latency_boxplot.gp
```

Example III




Figure – Waterfall visualization of data's journey

Paths found in /home/flavien/latseq.23042020.lseq

Downlink paths

path 0 : ip -> pdcp.in -> pdcp.tx -> rlc.tx.um -> rlc.seq.um -> mac.mux -> mac.t

path 1 : rlc.tx.am -> rlc.seq.am -> mac.mux -> mac.txreq -> phy.out.proc

Uplink paths

path 2 : phy.in.proc -> mac.demux -> rlc.rx.um -> rlc.unseq.um -> pdcp.rx -> ip

Stats for Journeys latency

Values Downlink

Size 51715

Average 2.58

StDev 1.07

Max 5.62

[75..90%] 3.88

[50..75%] 3.7

[25..50%] 2.81

[10..25%] 1.62

[0..10%] 0.996

Min 0.0341

Values Uplink

Size 900

Average 0.269

StDev 0.547

Max 2.04

[75..90%] 1.01

[50..75%] 0.0522

[25..50%] 0.01

[10..25%] 0.00691

[0..10%] 0.00596

Min 0.00477

Latency for points

Stats for Point Latency for mac.txreq

Values Downlink

Size 646

Average 0.0316

StDev 0.00356

Max 0.052

[75..90%] 0.0319

[50..75%] 0.031

[25..50%] 0.031

[10..25%] 0.03

[0..10%] 0.03

Min 0.03

Figure – Formatted latseq_stats output statistics

Scenarios

Scenarios to measure latency and Jitter. Do Statistic analysis...




Figure – Boxplot




Figure – cnorm




Figure – dots latency for timestamp

Conclusion

▶ What is already present

- ▶ A complete eNB NoS1, sufficient to run a trial of latseq
- ▶ PHY : Downlink for FR1, LDPC Decoder, polar coder
- ▶ MAC : FAPI P5
- ▶ RRC : all gNB messages are handled

▶ What is missing

- ▶ 5G NoS1 NSA !
- ▶ All the layer2 with 5G-NR changes

▶ what is coming

- ▶ PHY : PDCCH and PDSCH adds. Receiver PRACH, PUSCH, PUCCH, SRS
- ▶ MAC : RACH procedures, IR-HARQ, Dynamic MAC Scheduler, MAC-RLC interface
- ▶ updates PDCP and RLC to 5G
- ▶ F1 interface for split CU/DU
- ▶ (Xn interface for master eNB)

After code understanding and the Article [?] of 2019.

References I

[?][?][?][?][?]